

CSCI 2400

Introduction to Computer Systems

Fall 2014

**Professor Rick Han
Department of Computer Science
University of Colorado at Boulder**


**<http://www.cs.colorado.edu/~rhan>
rhan@cs.colorado.edu**

Goal of this course

- Understand how software executes on computer hardware

or...

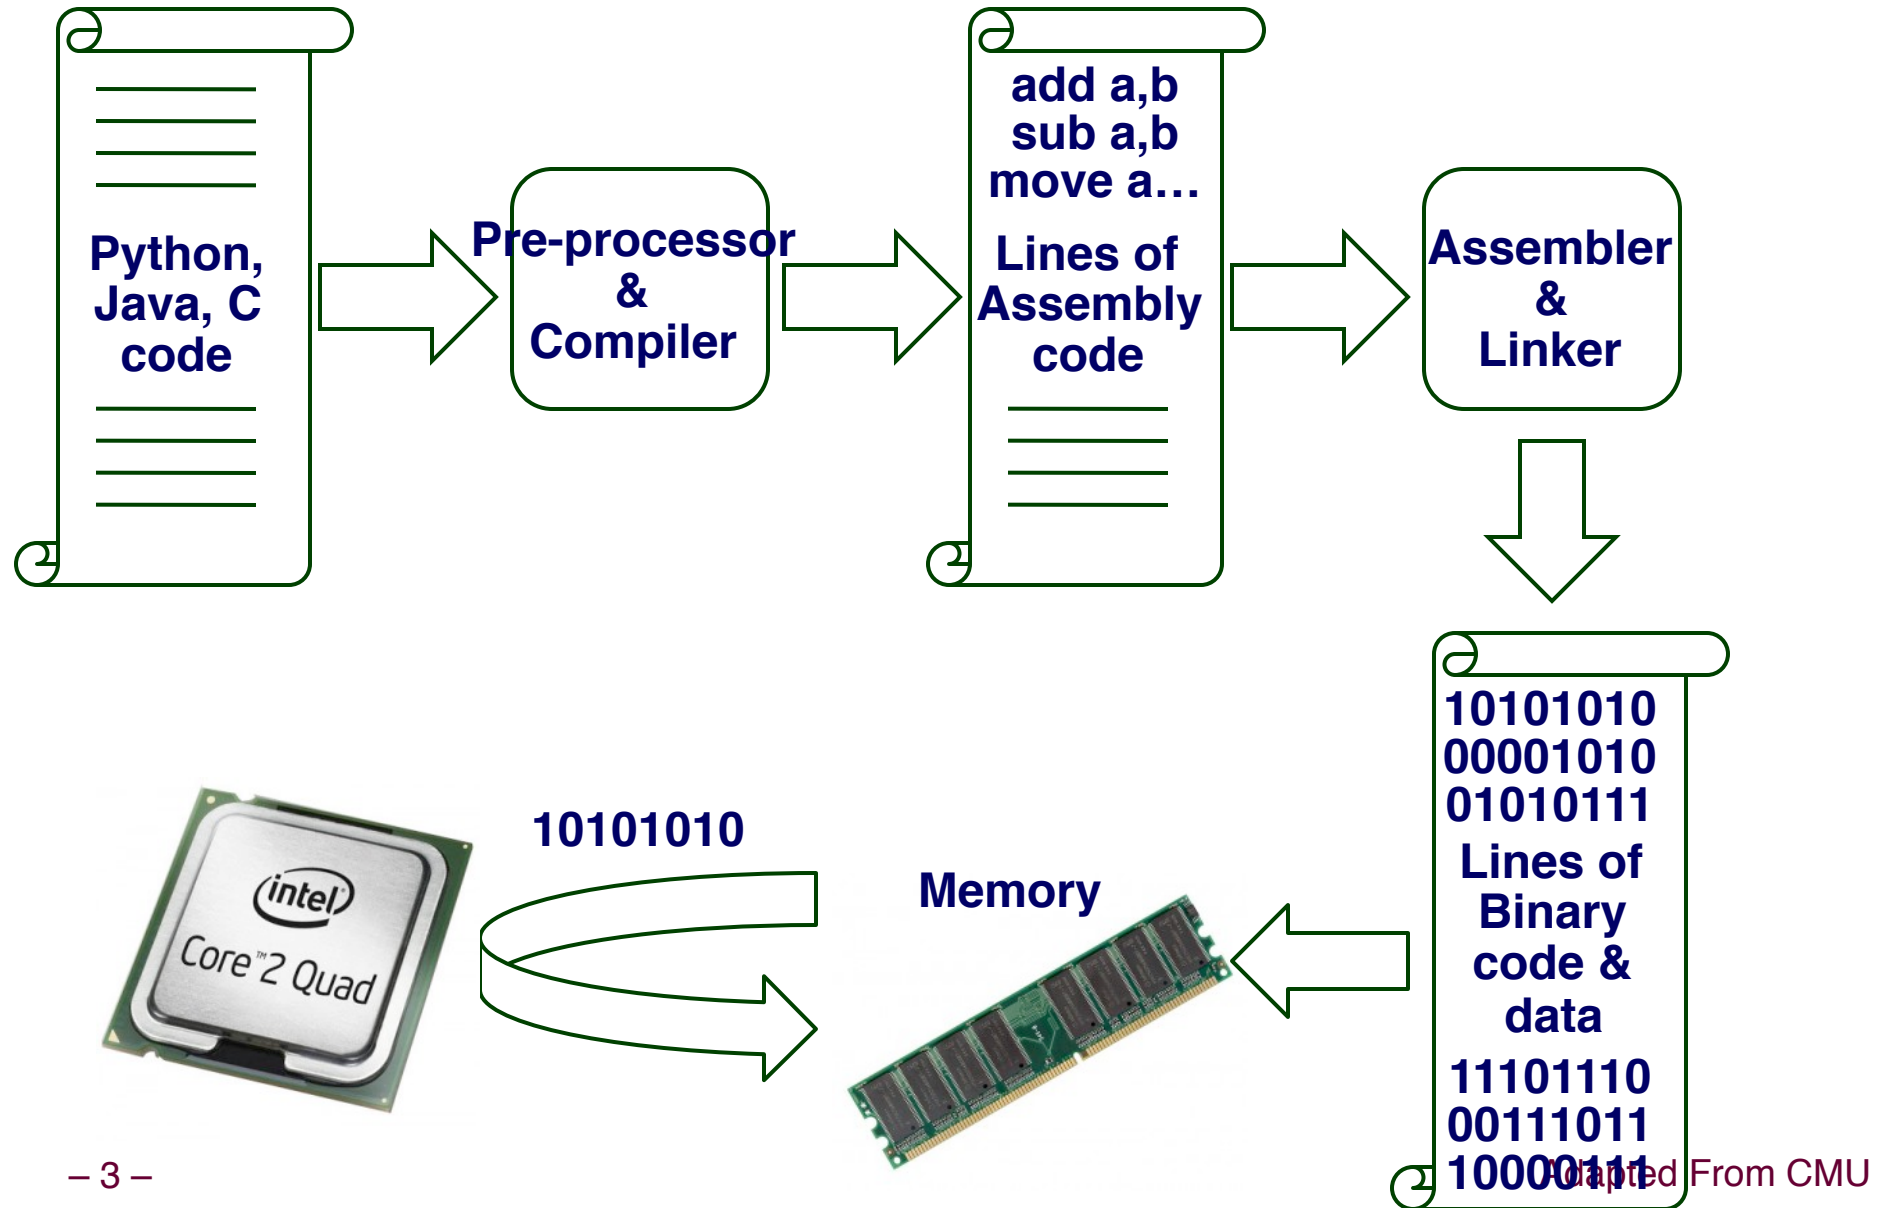
How a computer works!



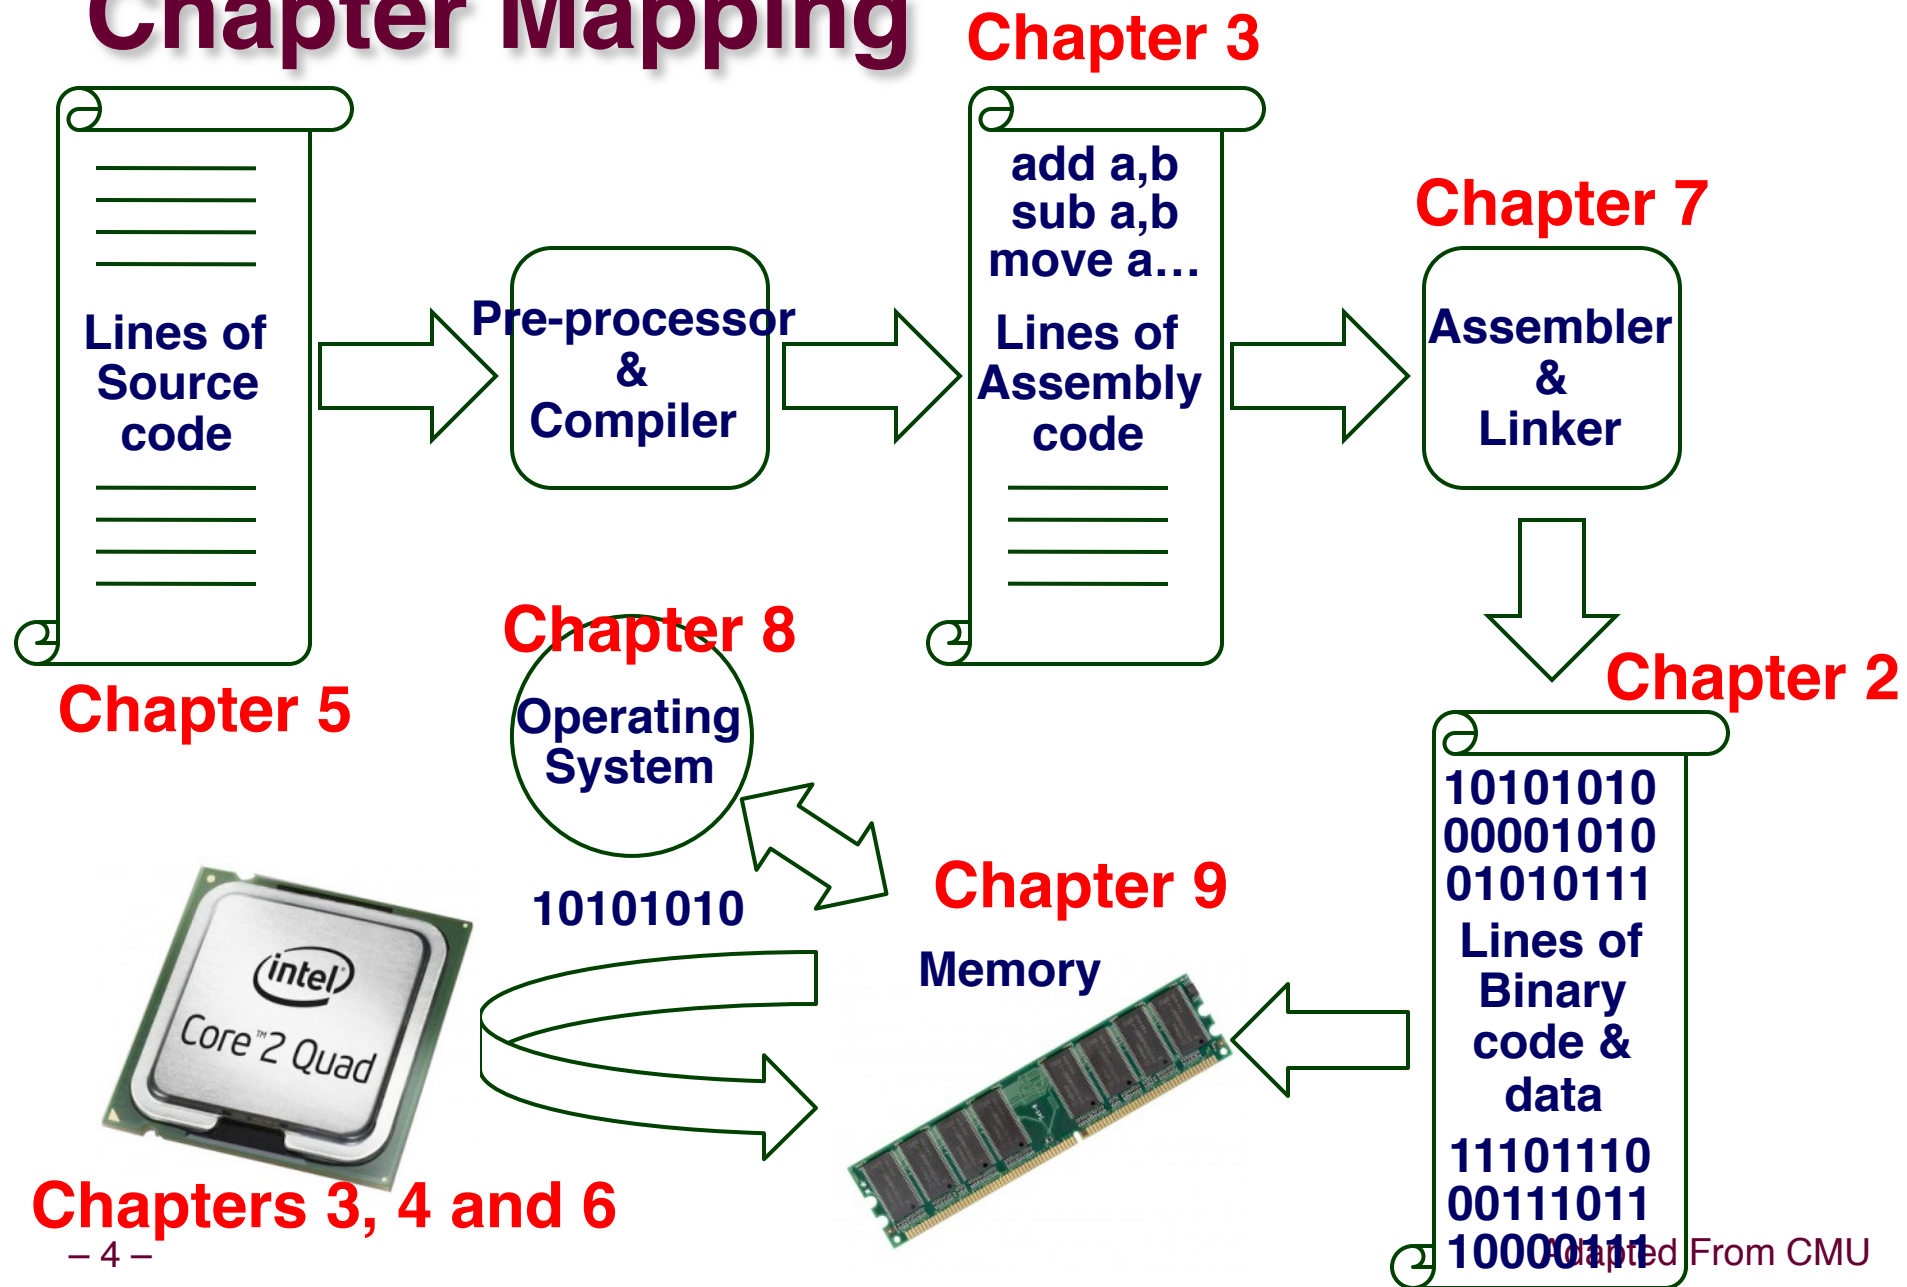
Great Material but Challenging!

Adapted From CMU

2400 In a Nutshell



Chapter Mapping



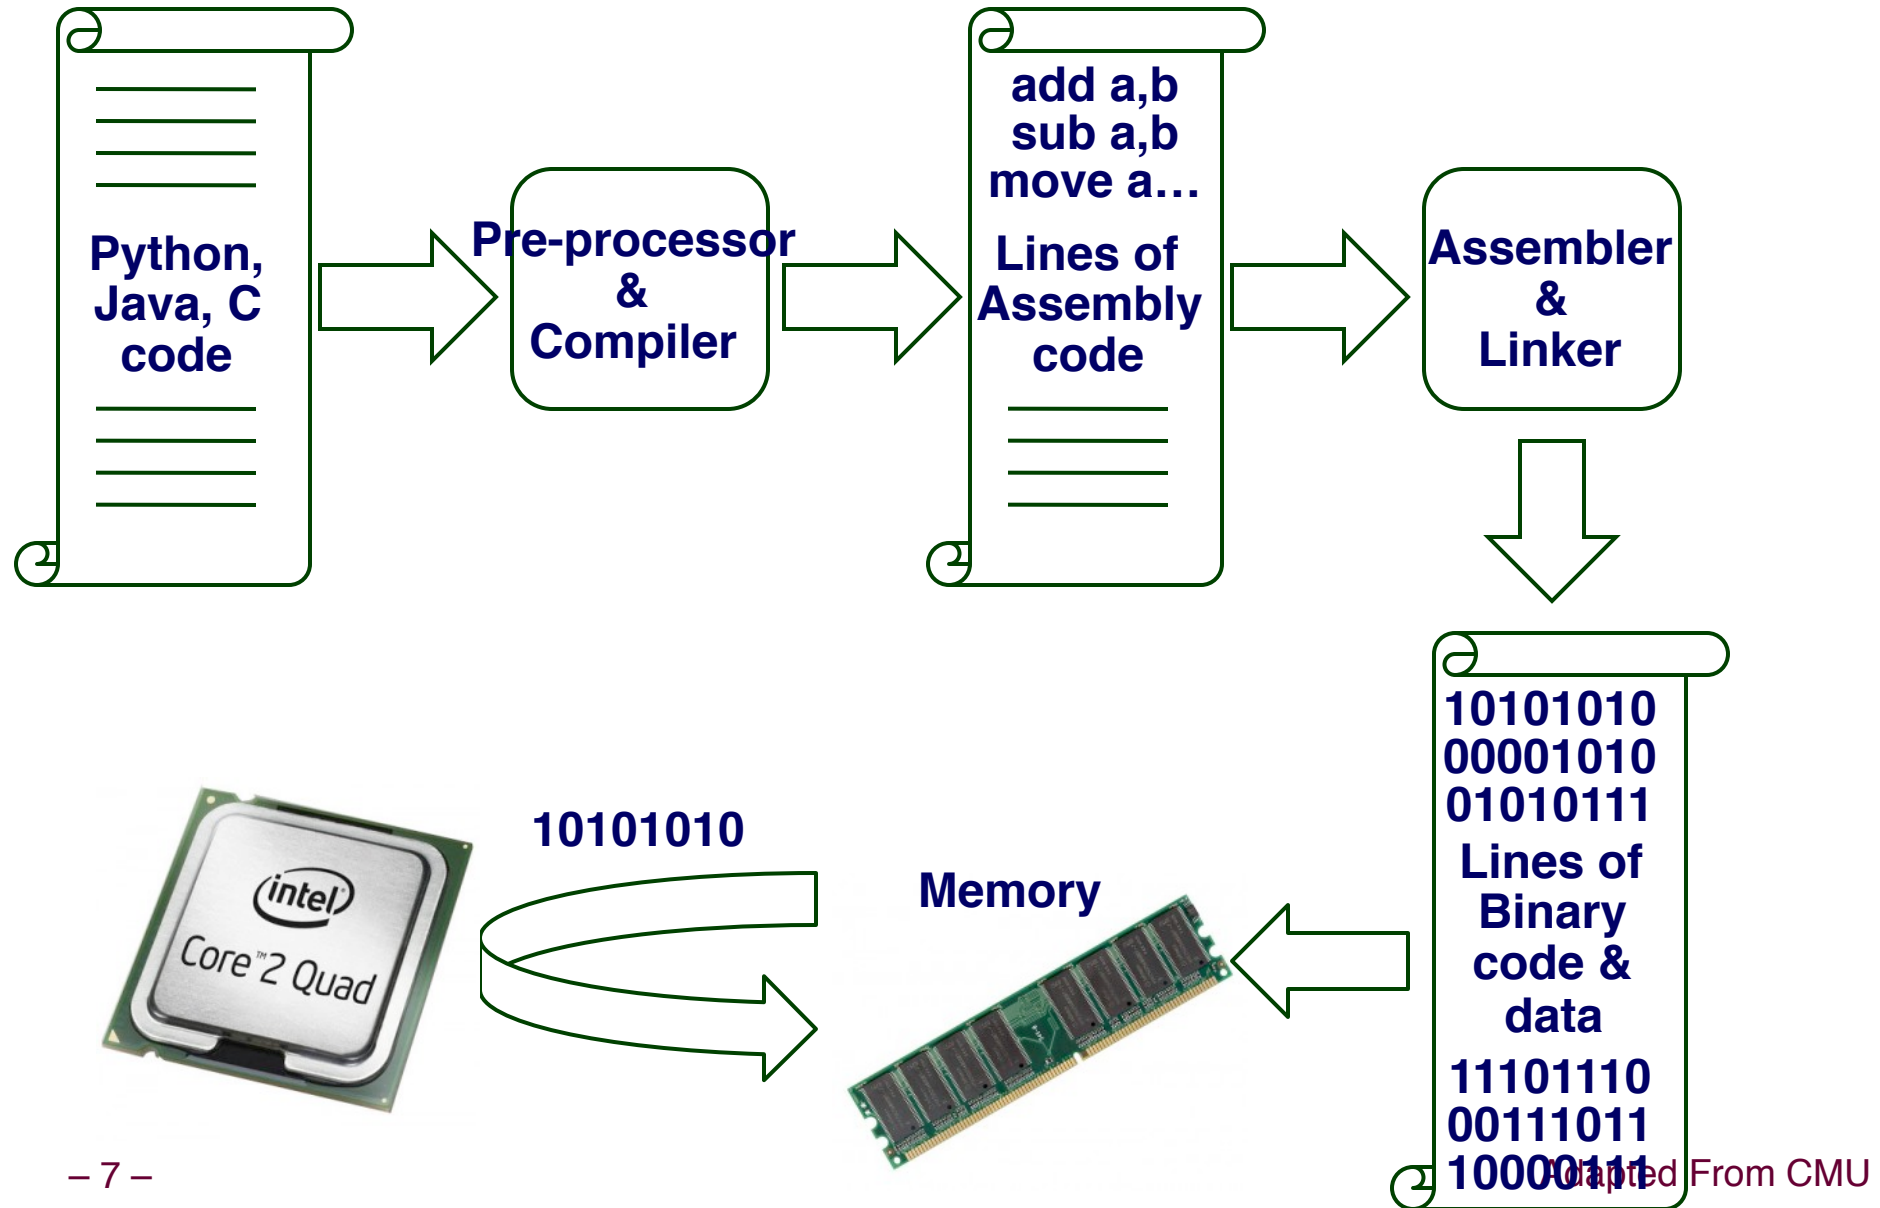
Approximate Timeline

Week	Topics	Due
1	Ch 1-2: Intro, Two's Complement	
2	Ch 2: Integer Arithmetic	
3	Ch 3: Assembly memory, arithmetic	Data Lab
4	Ch 3: Assembly control flow, loops	
5	Ch 3: Assembly: stacks, arrays	
6	Ch 3: Assembly: structs, 64-bit	Bomb Lab
7	Ch 2: Floating point	MT 1
8	Ch 4: ISA, Pipelining	
9	Ch 4-5: Pipelining, Optimization	Buffer Lab
10	Ch 5: Performance Optimization	
11	Ch 6: Caching	Performance Lab
12	Ch 6: Caching, Storage	MT 2
13	Ch 8: Exceptions, Signals	
14	Ch 9: Virtual Memory	Shell Lab
15	Ch 7: Linking	
Finals	Final Exam Thursday Dec 18, 4:30-7 pm	Final Exam

Announcements

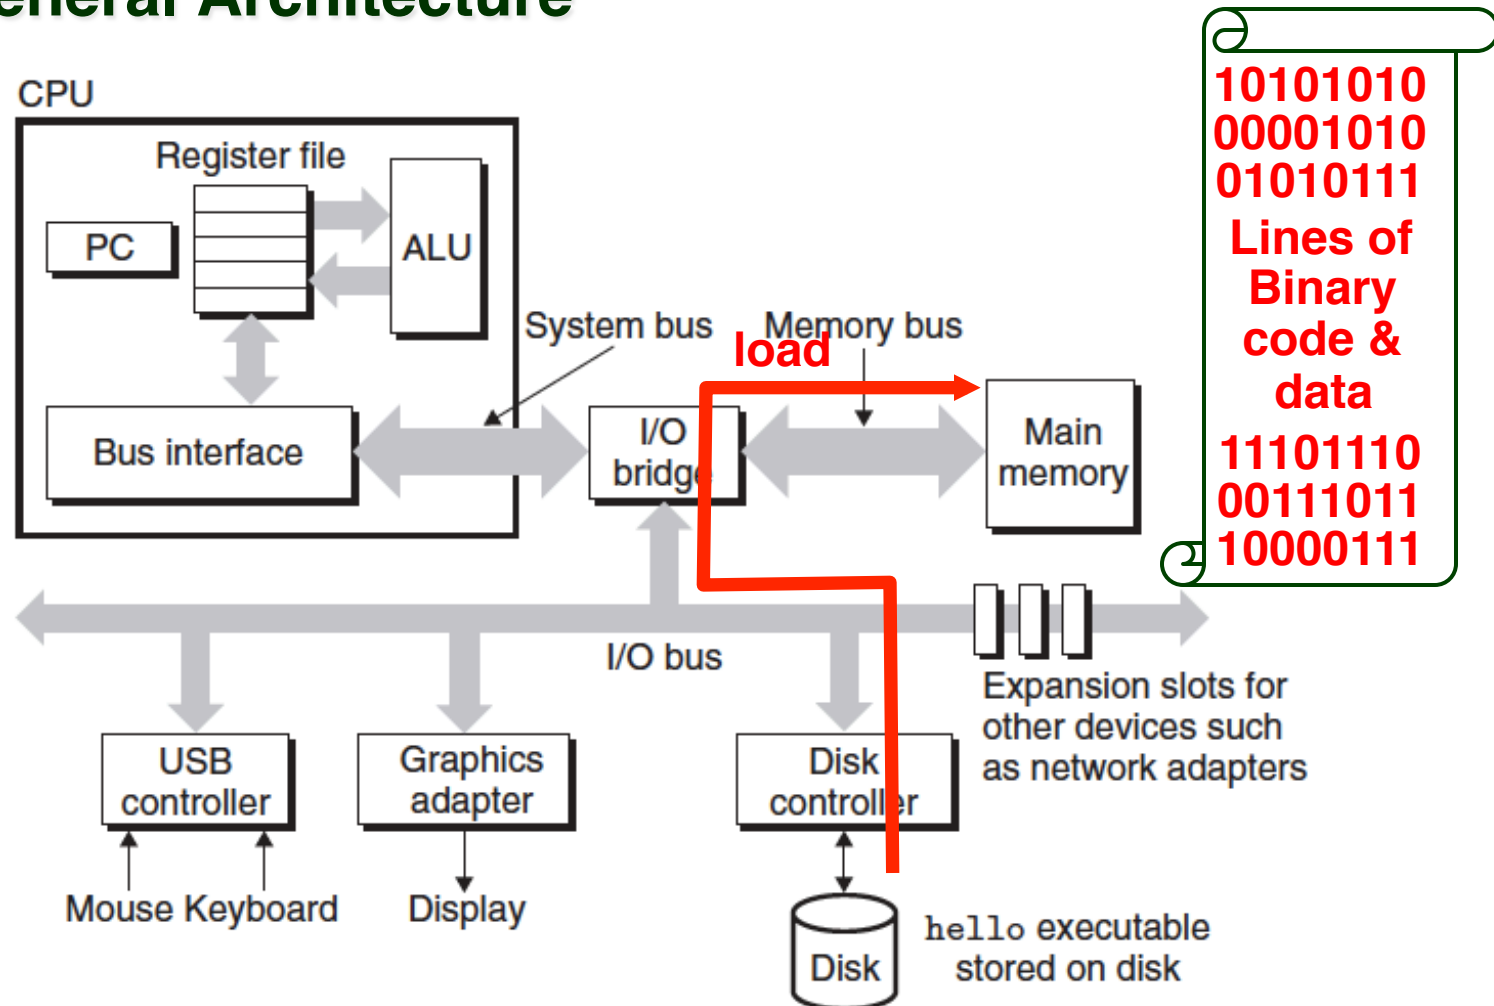
- **Let's go to moodle.cs.colorado.edu, the primary rendezvous point for CS 2400**
 - Sign up for an account, using '**systems14**' as the enrollment key
 - Walk through the syllabus
- **First data lab already released on moodle, due in ~3 weeks**
- **First recitation with TAs was on Monday**
 - Install the 2400 class VM (VirtualBox-based) on your laptops
 - More install sessions later this week – see link on moodle
- **Sign up for CSEL accounts @ csel.cs.colorado.edu**
- **Read Chapter 1 and 2.1-2.3 (skip floating point)**

2400 In a Nutshell



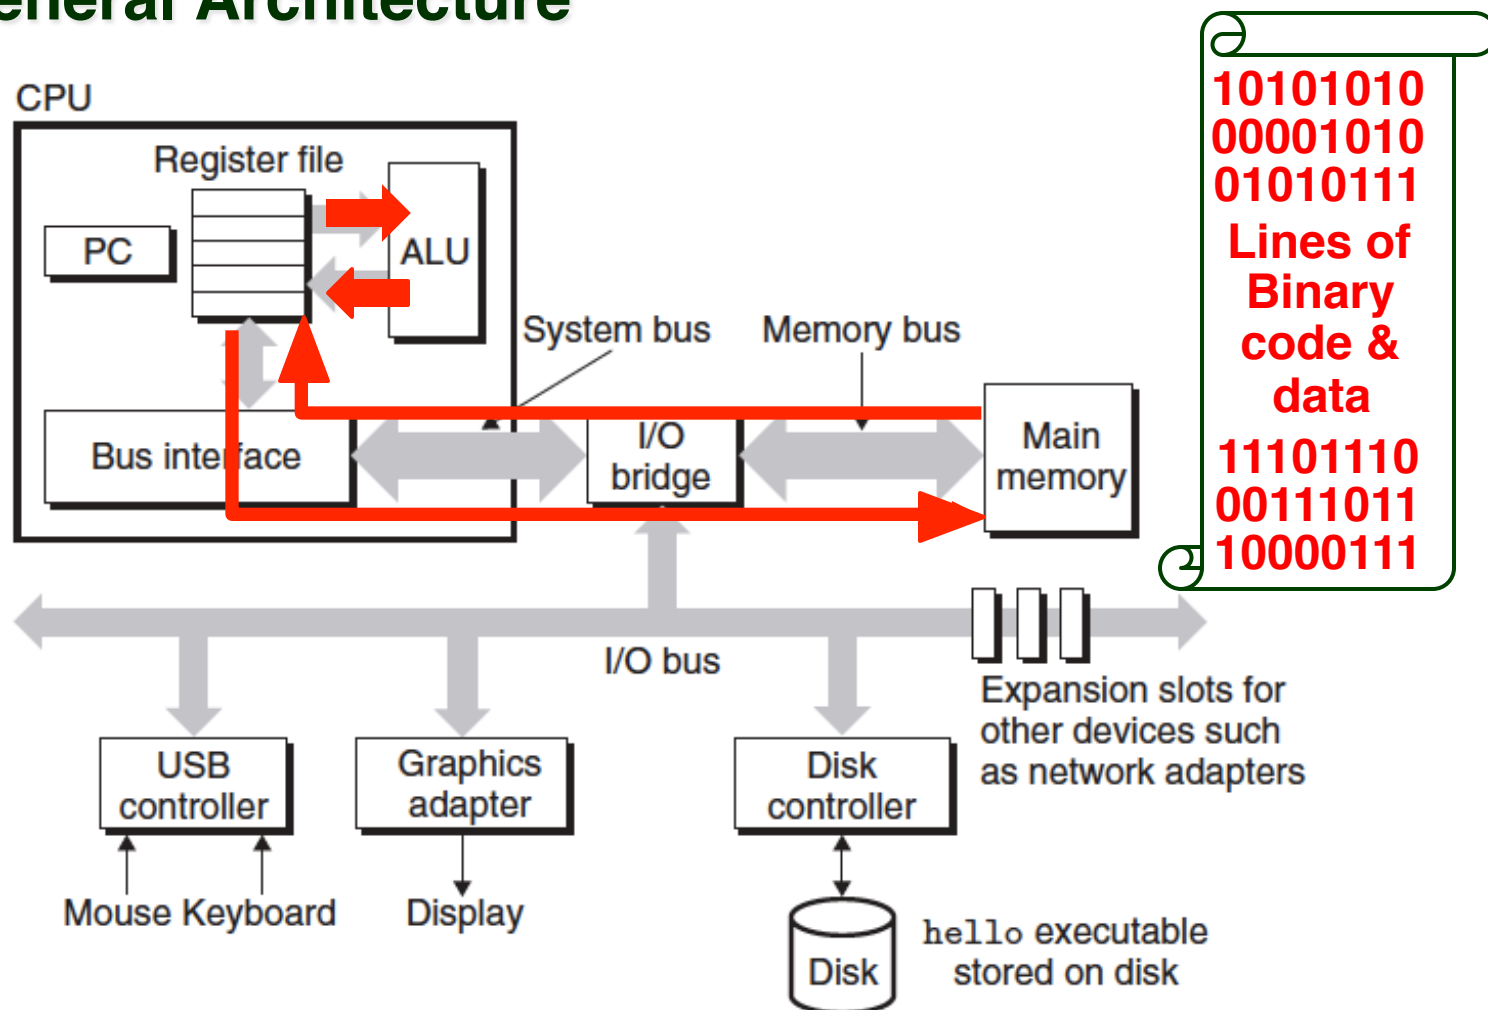
Hardware Organization of a Computer System (Von Neumann)

General Architecture




Hardware Organization of a Computer System

General Architecture



Hardware Organization of a Computer System

General Architecture




Inside the CPU - Building Blocks

- Registers

- Data (8 in IA32)
- Instruction register (IR) = current instruction
- Program counter (PC) = pointer to next instruction in memory

- Control Unit (not shown)

- Communicate with RAM




Arithmetic Logic Unit (ALU) of CPU

- Performs two types of operations

- Arithmetic: ADD, SUBTRACT, etc.
- Logic: AND, OR, NOT, XOR, >, <

- In the picture,

- A and B are the input data or operands
- F = function to perform, i.e. +, -, AND, XOR, etc.
- R = result
- D = any flags due to operation, such as a carry, overflow, sign, etc.



- A, B, and R are all stored in registers

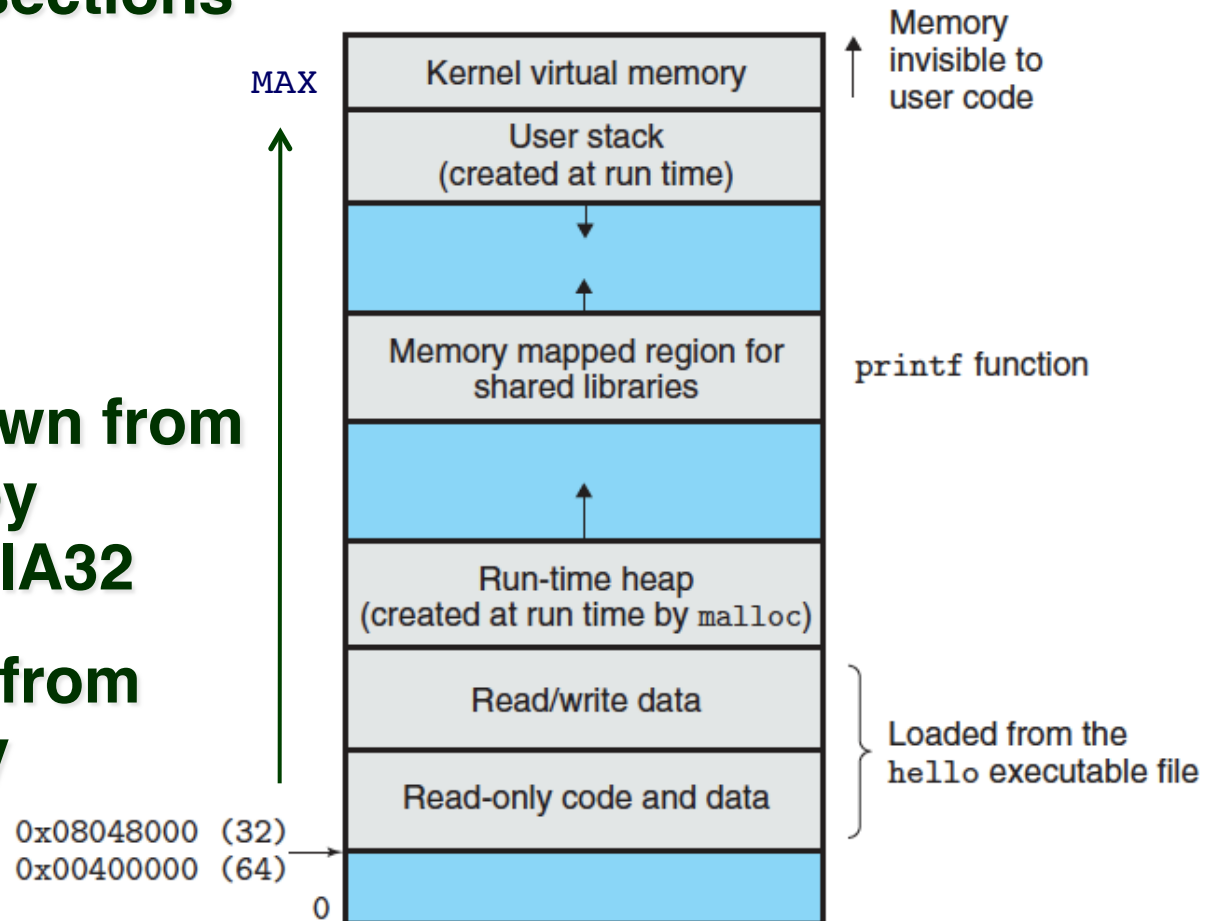
- As we will see, one type of instruction moves instructions and data to/from CPU from/to memory
 - i.e. fetch A and B from memory, move R to memory

Layout of a Program in Memory

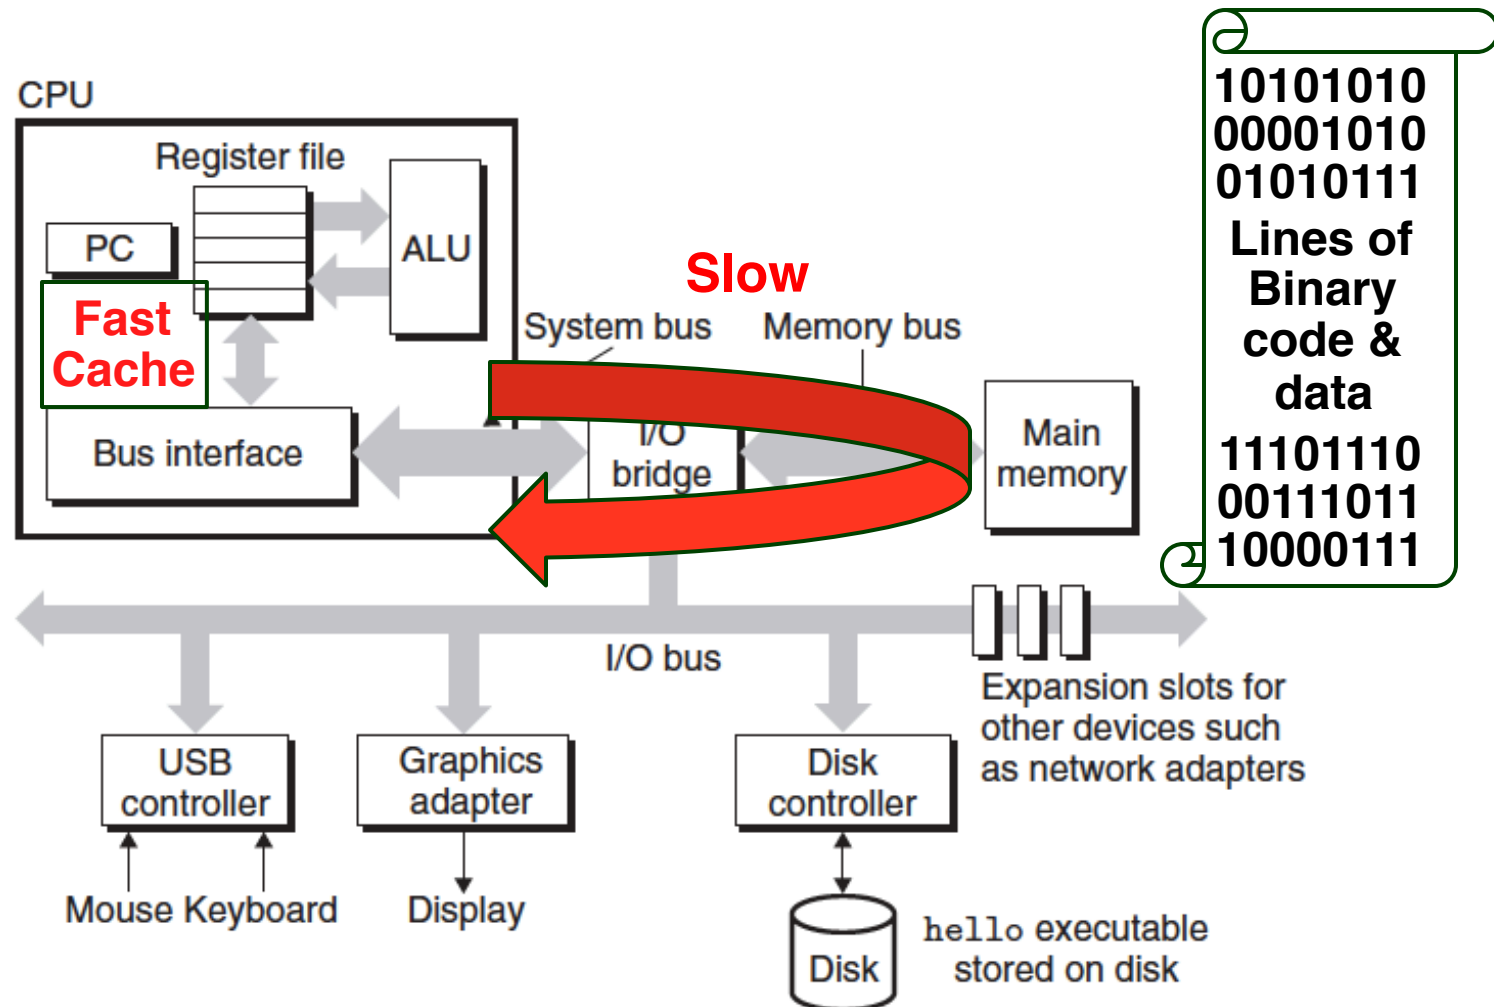
- Program memory is divided into 4 sections

1. Code
2. Data
3. Heap
4. Stack

- Stack grows down from high memory by convention on IA32
- Heap grows up from low memory by convention



Cache For Faster Data/Code Access

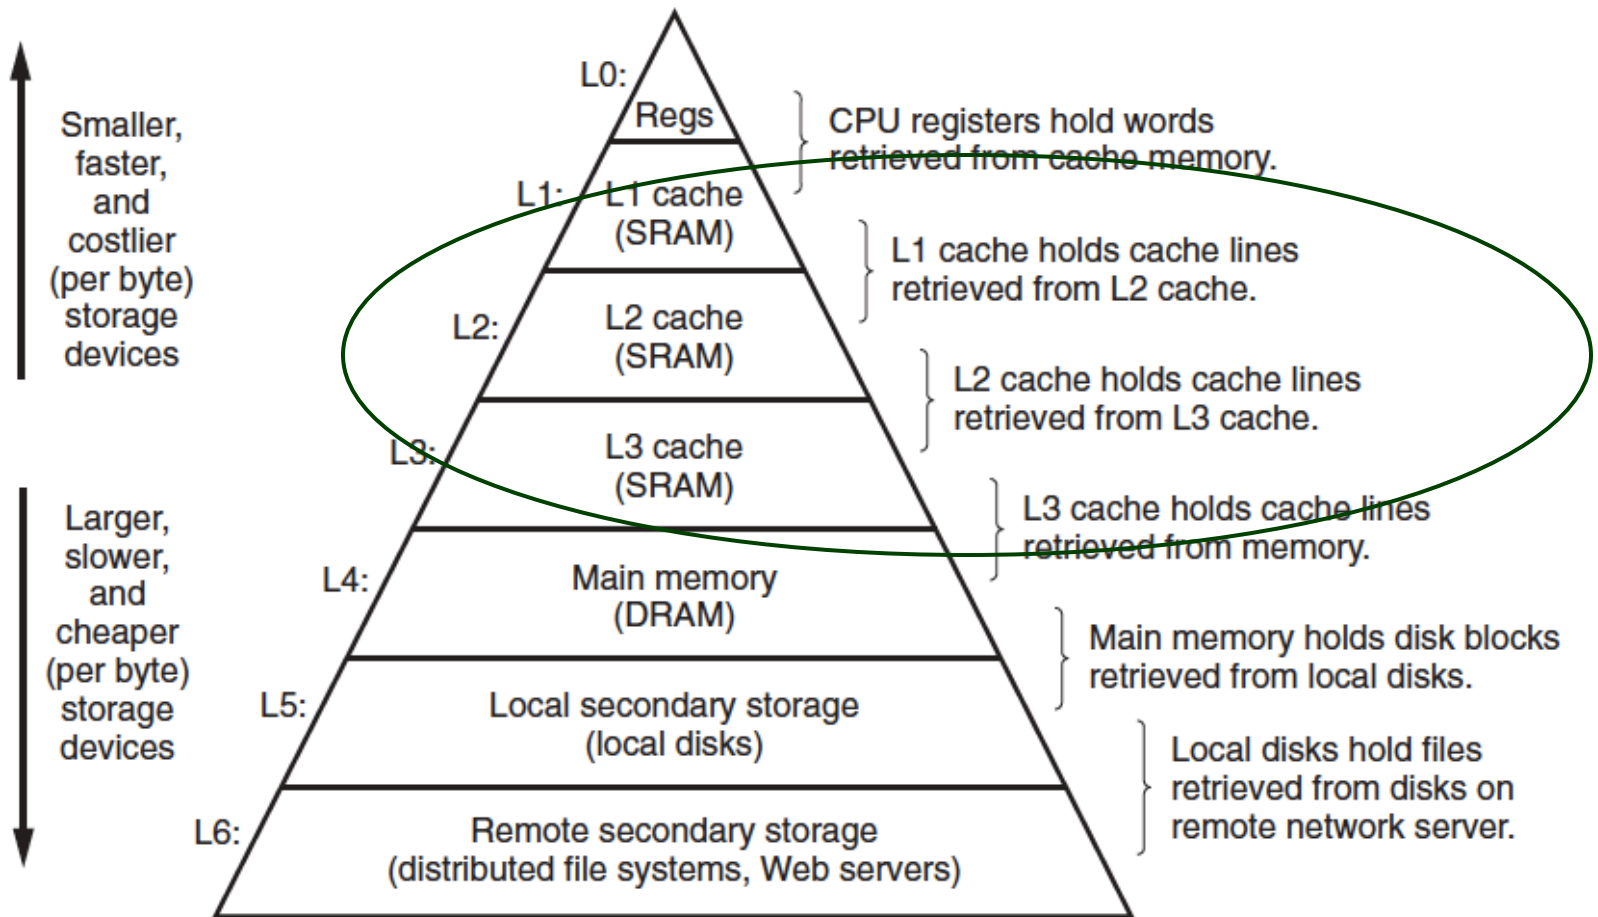


Cache For Faster Data/Code Access

▪ Caching


- Going to memory is quite slow compared to the speed of the CPU
 - RAM access time is ~ microseconds.
 - CPU executes multiple instructions per nanosecond.
 - So CPU waits ~1000 cycles to fetch data from memory!
- Solution: create a smaller but faster buffer called a cache, and cache data there that is going to be used soon
 - In reality, cache data that was more recently used
- Why not cache all data in registers?
 - Too expensive/byte

Memory Hierarchy



Memory Hierarchy

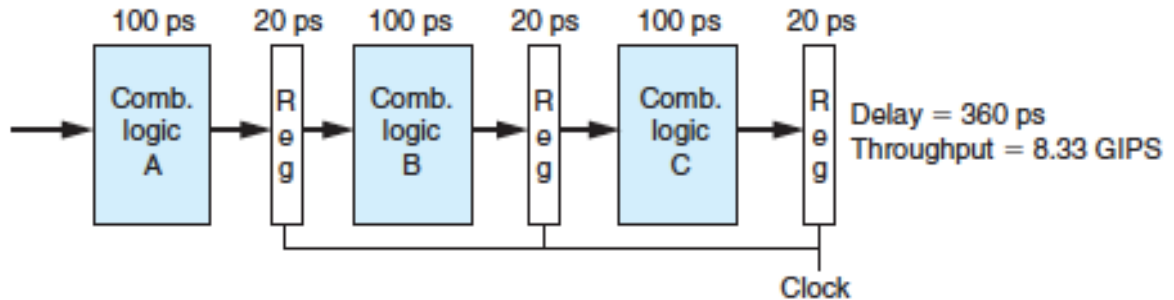
- **Intel Core i7 organization**
 - 4 cores
 - Each core has an L1 data cache, an L1 instruction cache, and a larger but slower L2 unified cache
 - Across cores, there is an L3 unified cache




Pipelining to Improve Throughput

■ Pipelining

- Instructions can be broken up into stages.
- Design CPU with multiple stages or “pipeline”.
- As a stage finishes, it accepts the result from the previous stage → Faster!
- Sequential non-pipelined is slower, some stages empty.



(a) Hardware: Three-stage pipeline



(b) Pipeline diagram

(a) Hardware: Unpipelined

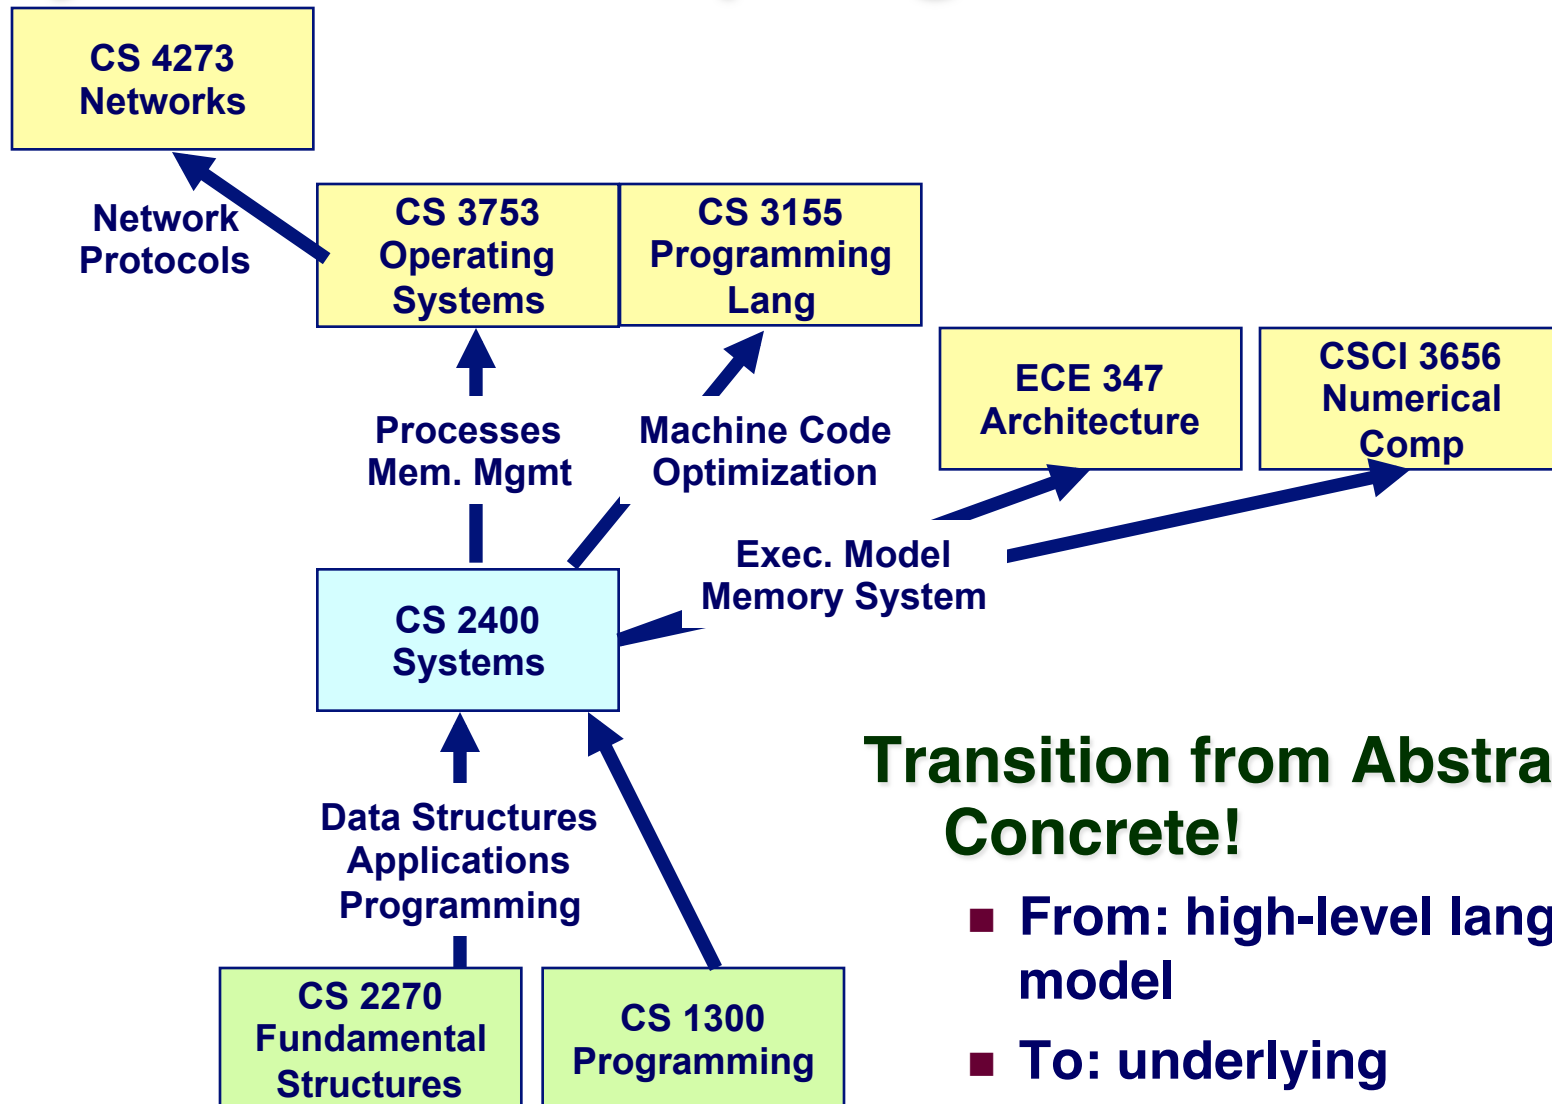


(b) Pipeline diagram

Why C?

- **A good compromise between a high-level programming language and low-level hardware**
 - **Maps well to assembly and binary machine code instructions**
 - **Low level enough to manipulate memory with pointers**
 - **Learn from C about the dangers of pointer arithmetic, array out-of-bounds memory accesses, etc. – we'll study some of these**
- **Most operating systems and network protocol stacks are built with C, as are many high-performance applications**

Systems as a Springboard...



Transition from Abstract to Concrete!

- From: high-level language model
- To: underlying implementation